

САНКТ-ПЕТЕРБУРГСКИЙ НАЦИОНАЛЬНЫЙ ИССЛЕДОВАТЕЛЬСКИЙ
УНИВЕРСИТЕТ ИТМО
Дисциплина: Архитектура ЭВМ

Отчет
по домашней работе №4
«ISA»

Выполнил(а): Альжанов Максим Булатович
студ. гр. М3139

Санкт-Петербург
2020

Теоретическая часть

ELF (executable and linkable format) – формат двоичных файлов часто используемый в unix системах. По своему дизайну ELF очень гибок, расширяем и кроссплатформенен. Например, он поддерживает возможность указывать порядок байтов или размеры адресов, чтобы не исключить возможность исполнения на некоторых и ISA.

Каждый ELF файл состоит из заголовка и данных. Заголовок состоит из такой информации как – класс (32 или 64 бита на адрес), ABI – описание интерфейса взаимодействия с операционной системой, целевая ISA, адрес entry – места откуда программа начнёт исполнение, адрес начала таблицы заголовков программ, адрес начала таблицы заголовков секций, индекс секции с именами секций в таблице заголовков секций и др.

Заголовок программы содержит всю необходимую информацию для размещения исполняемых данных в памяти компьютера.

Заголовок секции содержит в себе указатель на строку с названием секции, тип секции, фактическое расположение секции в данном файле. В секции могут храниться совершенно разные данные. Существует особая секция, которая содержит в себе строки с названиями секций. Её индекс хранится в заголовке elf файла.

RISC-V – открытая и свободная ISA основанная на концепции RISC. Основная ISA содержит в себе 53 команды, но может быть очень просто расширена. Существуют расширения для перемножения чисел (M), работы с плавающей точкой (F), сжатых команд (C), атомарных операций (A) и т.д.

RISC-V работает на 32 регистрах, соответственно для кодирования регистра нужно 5 бит. Регистры называют x%d, где %d – число от 0 до 31. X0 всегда равен нулю. По соглашению в X1 хранится указатель на возвращаемое

значение. В unix системах существуют соглашения по названию регистров (см. Таблица 1 - Соглашение о использовании регистров).

Таблица 1 - Соглашение о использовании регистров

Регистр	Название	Смысл
x0	zero	Ноль
x1	ra	Возвращаемое значение
x2	sp	Указатель на стек
x3	gp	Глобальный указатель
x4	tp	Указатель потока
x5-x7	t0-t2	Временные регистры
x8-x9	s0-s1	Регистры используемые вызывающим
x10-x17	a0-a7	Регистры аргументов
x18-x27	s2-s11	Регистры используемые вызывающим
x28-x31	t3-t6	Временные регистры

Базовая rv32i имеет длину инструкции 32 бита. Команды бывают нескольких типов – R, I, S, B, U и J. Каждая инструкция содержит opcode – располагается на семи младших битах. Opcode определяет длину инструкции для модификаций где длина инструкции не равна 32 битам.

R инструкция нужна для операций которые работают только на регистрах. Содержит 3 указателя на регистры: rs1, rs2, rd, два для чтения значений и один для записи, funct3, funct7 для определения операции.

I инструкция нужна для операций требующих временное значение imm (immediate) размером не больше 12 бит. Похожа на R, только место funct7 и rs2 занимает imm.

S инструкция нужна для записи значений в память. Похожа на R тип, но место rd и funct7 занимает imm – который в этих операциях играет роль дополнительного сдвиг для адреса памяти.

B инструкция нужна для условных переходов. Похожа на S тип, но imm записан по-другому.

U инструкция нужна для записи верхних бит 20 бит в какой либо регистр. Содержит только указатель на регистр и сохраняемое значение.

J инструкция нужна чтобы совершить прыжок в другое место.

Полное описание всех инструкций и соответствующие им opcode-ы можно найти в спецификации ISA.

Описание работы кода

- 1) Откроем файл для чтения
- 2) Прочитаем заголовок файла
- 3) Проверим что файл для RISC-V и для 32 бит.
- 4) Прочитаем таблицу секций
- 5) Найдём в ней секцию со строками, запомним, где она.
- 6) Найдём секцию .text.
- 7) Читаем по 4 байта и дисассемблируем каждую инструкцию. Если это инструкция прыжка и, то куда она прыгает не указывает на начало символа, то запишем адрес в множество «неизвестных адресов»
- 8) Пройдёмся по файлу ещё раз и будем снова дисассемблировать каждую инструкцию. Добавим метку в начало если у нас есть символ, указывающий на этот адрес или этот адрес есть во множестве «неизвестных адресов». Будем выводить построчно.
- 9) Закроем все файлы

Результат работы

```
00000000: <main>      addi sp, sp, -32
00000004:      sw ra, 28(sp)
00000008:      sw s0, 24(sp)
0000000C:      addi s0, sp, 32
00000010:      addi a0, zero, 0
00000014:      sw a0, 4084(s0)
00000018:      addi a1, zero, 64
0000001C:      sw a1, 4080(s0)
00000020:      sw a0, 4076(s0)
00000024:      addi a0, zero, 1
00000028:      sw a0, 4072(s0)
0000002C: <LOC_0x0000002C>  jal zero, 0 #0x0000002C <LOC_0x0000002C>
00000030:      lw a0, 4072(s0)
00000034:      lw a1, 4080(s0)
00000038: <LOC_0x00000038>  bge a0, a1, 0 #0x00000038 <LOC_0x00000038>
0000003C: <LOC_0x0000003C>  jal zero, 0 #0x0000003C <LOC_0x0000003C>
```

```

00000040:      lw a0, 4072(s0)
00000044:      mul a0, a0, a0
00000048:      lw a1, 4076(s0)
0000004C:      add a0, a1, a0
00000050:      sw a0, 4076(s0)
00000054: <LOC_0x00000054>      jal zero, 0 #0x00000054 <LOC_0x00000054>
00000058:      lw a0, 4072(s0)
0000005C:      addi a0, a0, 1
00000060:      sw a0, 4072(s0)
00000064: <LOC_0x00000064>      jal zero, 0 #0x00000064 <LOC_0x00000064>
00000068:      lw a0, 4076(s0)
0000006C:      lw s0, 24(sp)
00000070:      lw ra, 28(sp)
00000074:      addi sp, sp, 32
00000078:      jalr zero, ra, 0

```

Листинг кода

Язык: java 11 AdoptOpenJDK

src/Main.java

```

import me.alzhanov.ELF.RISCVDisassembler;
import net.fornwall.jelf.Elfile;

import java.io.*;

public class Main {
    public static void main(String[] args) {
        if (args.length < 1) {
            System.err.println("Usage: <input file> [<output file>]");
            return;
        }
        try {
            OutputStreamWriter output = null;
            try (BufferedInputStream stream = new BufferedInputStream(new
FileInputStream(args[0]))) {
                if (args.length > 1) {
                    output = new OutputStreamWriter(new
FileOutputStream(args[1]));
                } else {
                    output = new OutputStreamWriter(System.out);
                }
                RISCVDisassembler disassembler = new
RISCVDisassembler(Elfile.from(stream));
                disassembler.doDisassemble(new PrintWriter(output));
            } finally {
                if (output != null) {
                    output.close();
                }
            }
        } catch (FileNotFoundException e) {
            System.err.println("File is not found.");
        } catch (IOException e) {
            e.printStackTrace();
        }
    }
}

```

```
    }  
}
```

src/me/alzhanov/ELF/RISCVDisassembler.java

```
package me.alzhanov.ELF;  
  
import net.fornwall.jelf.*;  
  
import java.io.OutputStreamWriter;  
import java.io.PrintWriter;  
import java.util.HashSet;  
import java.util.InputMismatchException;  
import java.util.Set;  
import java.util.TreeSet;  
  
public class RISCVDisassembler {  
    final ElfFile file;  
  
    public RISCVDisassembler(ElfFile file) {  
        if (file.objectSize != ElfFile.CLASS_32) {  
            throw new InputMismatchException("That elf is not 32 bit.");  
        }  
        if (file.arch != 0xF3) {  
            throw new InputMismatchException("That elf is not for RISC-V.");  
        }  
        this.file = file;  
    }  
  
    public void dumpAll(OutputStreamWriter output) {  
        PrintWriter writer = new PrintWriter(output);  
        doDisassemble(writer);  
        dumpSymTable(writer);  
        writer.flush();  
    }  
  
    String getRegisterString(int reg) {  
        if (reg == 0)  
            return "zero";  
        else if (reg == 1)  
            return "ra";  
        else if (reg == 2)  
            return "sp";  
        else if (reg == 3)  
            return "gp";  
        else if (reg == 4)  
            return "tp";  
        else if (5 <= reg && reg <= 7)  
            return "t" + (reg - 5);  
        else if (reg == 8)  
            return "s0";  
        else if (reg == 9)  
            return "s1";  
        else if (10 <= reg && reg <= 17)  
            return "a" + (reg - 10);  
        else if (18 <= reg && reg <= 27)  
            return "s" + (reg - 18 + 2);  
        else if (28 <= reg && reg <= 31)  
            return "t" + (reg - 28 + 3);  
    }  
}
```

```

        else
            throw new AssertionError("RISC-V doesn't have register " + reg);
    }

    private String getSymbolForAddr(long loc, boolean isUnmarked) {
        ElfSymbol symb = file.getELFSymbol(loc);
        String locS = String.format("0x%08X", loc);
        if (symb != null && symb.st_value == loc && symb.section_type ==
ElfSymbol.STT_FUNC) {
            locS += " <" + symb.getName() + ">";
        } else if (isUnmarked) {
            locS += String.format(" <LOC_0x%08X>", loc);
        }
        return locS;
    }

    Set<Long> findUnmarkedLocations(ElfSection textSection) {
        long curOffset = 0;
        file.parser.seek(textSection.header.section_offset);
        Set<Long> syms = new HashSet<>();
        while (curOffset < textSection.header.size) {
            long virtualAddress = curOffset + textSection.header.address;
            int instruction = file.parser.readInt();
            int opcode = instruction & ((1 << 7) - 1);
            if (opcode == 0b1101111) { // JAL
                int offset = getOffsetForJType(instruction);
                ElfSymbol symb = file.getELFSymbol(virtualAddress + offset);
                if (symb == null || symb.st_value != virtualAddress + offset ||
symb.section_type != ElfSymbol.STT_FUNC)
                    syms.add(virtualAddress + offset);
            } else if (opcode == 0b1100011) { // B-type
                int offset = getOffsetForBType(instruction);
                ElfSymbol symb = file.getELFSymbol(virtualAddress + offset);
                if (symb == null || symb.st_value != virtualAddress + offset ||
symb.section_type != ElfSymbol.STT_FUNC)
                    syms.add(virtualAddress + offset);
            }
            curOffset += 4;
        }
        return syms;
    }

    private int getOffsetForBType(int instruction) {
        int offset = (((instruction >>> 8) & ((1 << 4) - 1)) << 1) |
            (((instruction >>> 25) & ((1 << 6) - 1)) << 5) |
            (((instruction >>> 7) & 1) << 11) |
            (((instruction >>> 31) & 1) << 12);
        if ((offset & (1 << 12)) != 0) {
            offset = -(-offset & ((1 << 12) - 1));
        }
        return offset;
    }

    private int getOffsetForJType(int instruction) {
        int imm = instruction >>> 12;
        int offset = (((imm >>> 9) & ((1 << 10) - 1)) << 1) |
            (((imm >>> 8) & 1) << 11) |
            ((imm & ((1 << 8) - 1)) << 12) |
            (((imm >>> 19) & 1) << 20);
    }

```

```

        if ((offset & (1 << 20)) != 0) {
            offset = -(-offset & ((1 << 20) - 1));
        }
        return offset;
    }

    public void doDisassemble(PrintWriter out) {
        file.getDynamicSymbolTableSection();
        file.getSymbolTableSection();
        ElfSection textSection = file.firstSectionByName(".text");
        if (textSection == null)
            throw new InputMismatchException("No .text found");
        long curOffset = 0;
        Set<Long> unmarked = findUnmarkedLocations(textSection);
        file.parser.seek(textSection.header.section_offset);
        while (curOffset < textSection.header.size) {
            long virtualAddress = curOffset + textSection.header.address;
            out.print(String.format("%08X: ", virtualAddress));
            int instruction = file.parser.readInt();
            ElfSymbol symb = file.getELFSymbol(virtualAddress);
            if (symb != null && symb.st_value == virtualAddress &&
                symb.section_type == ElfSymbol.STT_FUNC) {
                out.printf("<%s>\t", symb.getName());
            } else if (unmarked.contains(virtualAddress)) {
                out.printf("<LOC_0x%08X>\t", virtualAddress);
            } else {
                out.print("\t");
            }
            int opcode = instruction & ((1 << 7) - 1);
            int rd = instruction >> 7 & ((1 << 5) - 1);
            int funct3 = instruction >> 12 & ((1 << 3) - 1);
            int rs1 = instruction >> 15 & ((1 << 5) - 1);
            int rs2 = instruction >> 20 & ((1 << 5) - 1);
            int imm110 = instruction >> 20 & ((1 << 12) - 1);
            int funct7 = instruction >> 25;
            if (opcode == 0b0110111) { // LUI
                out.printf("%6s %s, %s\n", "lui", getRegisterString(rd),
                    Integer.toUnsignedString((instruction >>> 12) << 12));
            } else if (opcode == 0b0010111) { // AUIPC
                out.printf("%6s %s, %s\n", "auipc", getRegisterString(rd),
                    Integer.toUnsignedString((instruction >>> 12) << 12));
            } else if (opcode == 0b110111) { // JAL
                int offset = getOffsetForJType(instruction);
                long jumpTo = virtualAddress + offset;
                out.printf("%6s %s, %d\t#%s\n", "jal", getRegisterString(rd),
                    offset, getSymbolForAddr(jumpTo, unmarked.contains(jumpTo)));
            } else if (opcode == 0b1100111 && funct3 == 0b000) { // jalr
                if ((imm110 & (1 << 11)) != 0) {
                    imm110 = -(-imm110 & ((1 << 11) - 1));
                }
                out.printf("%6s %s, %s, %d\n", "jalr", getRegisterString(rd),
                    getRegisterString(rs1), imm110);
            } else if (opcode == 0b1100011) { // B-type
                int offset = getOffsetForBType(instruction);
                String instr = new String[]{"beq", "bne", "??", "??", "blt",
                    "bge", "bltu", "bgeu"}[funct3];
                long jumpTo = virtualAddress + offset;
                out.printf("%6s %s, %s, %d\t#%s %n", instr,
                    getRegisterString(rs1), getRegisterString(rs2), offset, getSymbolForAddr(jumpTo,

```



```

unmarked.contains(jumpTo)));
    } else if (opcode == 0b0000011) { // I-type - LB, LH, LW, LBU, LHU
        String instr = new String[]{"lb", "lh", "lw", "??", "lbu", "lhu",
"??", "??"}[funct3];
        out.printf("%6s %s, %d(%s)%n", instr, getRegisterString(rd),
imm110, getRegisterString(rs1));
    } else if (opcode == 0b0100011) { // S-type SB, SH, SW
        String instr = new String[]{"sb", "sh", "sw", "??", "??", "??",
"??", "??"}[funct3];
        int imm = rd | ((imm110 >> 5) << 5);
        out.printf("%6s %s, %d(%s)%n", instr, getRegisterString(rs2), imm,
getRegisterString(rs1));
    } else if (opcode == 0b0010011) {
        if (funct3 == 0b001) { // SLLI
            out.printf("%6s %s, %s, %d%n", "slli", getRegisterString(rd),
getRegisterString(rs1), imm110);
        } else if (funct3 == 0b101) {
            if (funct7 == 0b0100000) { // SRAI
                out.printf("%6s %s, %s, %d%n", "srai",
getRegisterString(rd), getRegisterString(rs1), imm110 & ((1 << 5) - 1));
            } else { // SRLI
                out.printf("%6s %s, %s, %d%n", "srli",
getRegisterString(rd), getRegisterString(rs1), imm110);
            }
        } else { // I-type - ADDI, SLTI, SLTIU, XORI, ORI, ANDI
            String instr = new String[]{"addi", "??", "slti", "sltiu",
"xori", "??", "ori", "andi"}[funct3];
            if (instr.equals("addi") || instr.equals("slti")) { // sign-
extend
                if ((imm110 & (1 << 11)) != 0) {
                    imm110 = -(-imm110 & ((1 << 11) - 1));
                }
            }
            out.printf("%6s %s, %s, %d%n", instr, getRegisterString(rd),
getRegisterString(rs1), imm110);
        }
    } else if (opcode == 0b110011) { // R-type
        if (funct7 == 0b0100000) { // SUB, SRA
            String instr = new String[]{"sub", "??", "??", "??", "??",
"sra", "??", "??"}[funct3];
            out.printf("%6s %s, %s, %s%n", instr, getRegisterString(rd),
getRegisterString(rs1), getRegisterString(rs2));
        } else if (funct7 == 0) {
            String instr = new String[]{"add", "sll", "slt", "sltu",
"xor", "srl", "or", "and"}[funct3];
            out.printf("%6s %s, %s, %s%n", instr, getRegisterString(rd),
getRegisterString(rs1), getRegisterString(rs2));
        } else if (funct7 == 1) {
            String instr = new String[]{"mul", "mulh", "mulhsu", "mulhu",
"div", "divu", "rem", "remu"}[funct3];
            out.printf("%6s %s, %s, %s%n", instr, getRegisterString(rd),
getRegisterString(rs1), getRegisterString(rs2));
        }
    } else if (opcode == 0b0001111) {
        if (funct3 == 1) { // FENCE.I
            out.printf("%6s%n", "fence.i");
        } else { // FENCE
            out.printf("%6s %d, %d%n", "fence", imm110 >> 4 << 4, imm110
& ((1 << 4) - 1));
        }
    }
}

```

```

    }
    } else if (opcode == 0b1110011) {
        if (funct3 == 0) {
            if (imm110 == 0) { // ECALL
                out.printf("%6s%n", "ecall");
            } else if (imm110 == 1) { // EBREAK
                out.printf("%6s%n", "ebreak");
            } else {
                out.printf("????%n");
            }
        } else {
            String instr = new String[]{"", "csrrw", "csrrs", "csrrc",
"??", "csrrwi", "csrrsi", "csrrci"}[funct3];
            out.printf("%6s %s, %s, %s%n", instr, getRegisterString(rd),
imm110, getRegisterString(rs1));
        }
    } else {
        out.printf("????%n");
    }
    curOffset += 4;
}

private static int getIntWidth(int a) {
    if (a == 0)
        return 1;
    return (int) Math.floor(Math.log10(Math.abs(a))) + 1 + (a < 0 ? 1 : 0);
}

static String symbolTypeToString(int type) {
    switch (type) {
        case (ElfSymbol.STT_NOTYPE):
            return "NOTYPE";
        case (ElfSymbol.STT_OBJECT):
            return "OBJECT";
        case (ElfSymbol.STT_FUNC):
            return "FUNC";
        case (ElfSymbol.STT_SECTION):
            return "SECTION";
        case (ElfSymbol.STT_FILE):
            return "FILE";
        case (ElfSymbol.STT_LOPROC):
            return "LOPROC";
        case (ElfSymbol.STT_HIPROC):
            return "HIPROC";
        default:
            return "UNKNOWN";
    }
}

static String bindingToString(int binding) {
    switch (binding) {
        case (ElfSymbol.BINDING_GLOBAL):
            return "GLOBAL";
        case (ElfSymbol.BINDING_HIPROC):
            return "HIPROC";
        case (ElfSymbol.BINDING_LOCAL):
            return "LOCAL";
        case (ElfSymbol.BINDING_LOPROC):

```

```

        return "LOPROC";
    case (ElfSymbol.BINDING_WEAK):
        return "WEAK";
    default:
        return "UNKNOWN";
    }
}

static String visibilityToString(ElfSymbol.Visibility visibility) {
    switch (visibility) {
        case STV_HIDDEN:
            return "HIDDEN";
        case STV_DEFAULT:
            return "DEFAULT";
        case STV_INTERNAL:
            return "INTERNAL";
        case STV_PROTECTED:
            return "PROTECTED";
        default:
            return "UNKNOWN";
    }
}

public void dumpSymTable(PrintWriter out) {
    out.println("Symtable:");
    ElfSymbolTableSection symtable = file.getSymbolTableSection();
    int symbolCount = symtable.symbols.length;
    int firstColWidth = getIntWidth(symbolCount);
    out.println(String.format("%" + (firstColWidth + 2) + "s   %8s %5s %7s %7s
%8s %4s %s",
        "Symbol".substring(0, firstColWidth + 2), "Value", "Size", "Type",
"Bind", "Vis", "Index", "Name"));
    for (int i = 0; i < symbolCount; i++) {
        ElfSymbol symbol = symtable.symbols[i];
        out.println(String.format("[%"+ firstColWidth + "s] 0x%08X %5s %7s
%7s %8s %4s %s",
            i,
            symbol.st_value,
            symbol.st_size,
            symbolTypeToString(symbol.getType()),
            bindingToString(symbol.getBinding()),
            visibilityToString(symbol.getVisibility()),
            shindexToString(symbol.st_shndx),
            symbol.st_name == 0 ? "" : symbol.getName()
        ));
    }
}

private String shindexToString(short stShndx) {
    if (stShndx == ElfSectionHeader.SHN_ABS) {
        return "ABS";
    } else if (stShndx == ElfSectionHeader.SHN_COMMON) {
        return "COMMON";
    } else if (Short.compareUnsigned(ElfSectionHeader.SHN_LOPROC, stShndx) <=
0 && Short.compareUnsigned(stShndx, ElfSectionHeader.SHN_HIPROC) <= 0) {
        return "PROC_RES";
    } else if (Short.compareUnsigned(ElfSectionHeader.SHN_LOOS, stShndx) <= 0
&& Short.compareUnsigned(stShndx, ElfSectionHeader.SHN_HIOS) <= 0) {
        return "OS_RES";
    }
}

```

```

        } else if (stShndx == ElfSectionHeader.SHN_UNDEF) {
            return "UNDEF";
        } else if (stShndx == ElfSectionHeader.SHN_XINDEX) {
            return "XINDEX";
        } else if (ElfSectionHeader.SHN_LORESERVE <= stShndx && stShndx <=
ElfSectionHeader.SHN_HIRESERVE) {
            return "RESERVED";
        } else {
            return String.valueOf(stShndx);
        }
    }
}

```

src/net/fornwall/jelf/BackingFile.java

```

package net.fornwall.jelf;

import java.io.ByteArrayInputStream;
import java.io.IOException;
import java.nio.Buffer;
import java.nio.MappedByteBuffer;

class BackingFile {
    private final ByteArrayInputStream byteArray;
    private final MappedByteBuffer mappedByteBuffer;
    private final long mbbStartPosition;

    public BackingFile(ByteArrayInputStream byteArray) {
        this.byteArray = byteArray;
        this.mappedByteBuffer = null;
        this.mbbStartPosition = -1;
    }

    public BackingFile(MappedByteBuffer mappedByteBuffer) {
        this.byteArray = null;
        this.mappedByteBuffer = mappedByteBuffer;
        this.mbbStartPosition = 0;
        ((Buffer)mappedByteBuffer).position((int) mbbStartPosition);
    }

    public void seek(long offset) {
        if (byteArray != null) {
            byteArray.reset();
            if (byteArray.skip(offset) != offset) throw new ElfException("seeking
outside file");
        } else if (mappedByteBuffer != null) {
            ((Buffer)mappedByteBuffer).position((int)(mbbStartPosition + offset));
            // we may be limited to sub-4GB mapped files
        }
    }

    public void skip(int bytesToSkip) {
        if (byteArray != null) {
            long skipped = byteArray.skip(bytesToSkip);
            if (skipped != bytesToSkip) {
                throw new IllegalArgumentException("Wanted to skip " + bytesToSkip
+ " bytes, but only able to skip " + skipped);
            }
        } else {

```

```

        ((Buffer)mappedByteBuffer).position(mappedByteBuffer.position() +
bytesToSkip);
    }
}

short readUnsignedByte() {
    int val = -1;
    if (byteArray != null) {
        val = byteArray.read();
    } else if (mappedByteBuffer != null) {
        byte temp = mappedByteBuffer.get();
        val = temp & 0xFF; // bytes are signed in Java == so assigning them
to a longer type risks sign extension.
    }

    if (val < 0) throw new ElfException("Trying to read outside file");
    return (short) val;
}

public int read(byte[] data) {
    if (byteArray != null) {
        try {
            return byteArray.read(data);
        } catch (IOException e) {
            throw new RuntimeException("Error reading " + data.length + "
bytes", e);
        }
    } else if (mappedByteBuffer != null) {
        mappedByteBuffer.get(data);
        return data.length;
    }
    throw new RuntimeException("No way to read from file or buffer");
}
}

```

src/net/fornwall/jelf/ElfDynamicSection.java

```
package net.fornwall.jelf;
```

```
import java.util.ArrayList;
import java.util.List;
```

```
/**
 * An {@link ElfSection} with information necessary for dynamic linking.
 * <p>
 * Given an {@link ElfFile}, use {@link ElfFile#getDynamicSection()} to obtain the
dynamic section for it if one exists,
 * which it only does if the ELF file is an object file participating in dynamic
linking.
 * <p>
 * This dynamic linking section contains a list of {@link ElfDynamicStructure}s.
 * <pre>
 * Name                Value  d_un          Executable  Shared Object
 * -----
 * DT_NULL              0    ignored      mandatory   mandatory
 * DT_NEEDED            1    d_val       optional    optional
 * DT_PLTRELSZ          2    d_val       optional    optional
 * DT_PLTGOT            3    d_ptr       optional    optional

```

```

* DT_HASH                4 d_ptr mandatory mandatory
* DT_STRTAB              5 d_ptr mandatory mandatory
* DT_SYMTAB              6 d_ptr mandatory mandatory
* DT_RELA                7 d_ptr mandatory optional
* DT_RELASZ              8 d_val mandatory optional
* DT_RELAENT             9 d_val mandatory optional
* DT_STRSZ              10 d_val mandatory mandatory
* DT_SYMENT             11 d_val mandatory mandatory
* DT_INIT               12 d_ptr optional optional
* DT_FINI               13 d_ptr optional optional
* DT_SONAME             14 d_val ignored optional
* DT_RPATH*            15 d_val optional ignored
* DT_SYMBOLIC*          16 ignored ignored optional
* DT_REL                17 d_ptr mandatory optional
* DT_RELSZ              18 d_val mandatory optional
* DT_RELENT             19 d_val mandatory optional
* DT_PLTREL             20 d_val optional optional
* DT_DEBUG              21 d_ptr optional ignored
* DT_TEXTREL*           22 ignored optional optional
* DT_JMPREL             23 d_ptr optional optional
* DT_BIND_NOW*          24 ignored optional optional
* DT_INIT_ARRAY         25 d_ptr optional optional
* DT_FINI_ARRAY         26 d_ptr optional optional
* DT_INIT_ARRAYSZ       27 d_val optional optional
* DT_FINI_ARRAYSZ       28 d_val optional optional
* DT_RUNPATH            29 d_val optional optional
* DT_FLAGS               30 d_val optional optional
* DT_ENCODING           32 unspecified unspecified unspecified
* DT_PREINIT_ARRAY      32 d_ptr optional ignored
* DT_PREINIT_ARRAYSZ    33 d_val optional ignored
* DT_LOOS                0x6000000D unspecified unspecified unspecified
* DT_HIOS                0x6ffff000 unspecified unspecified unspecified
* DT_LOPROC              0x70000000 unspecified unspecified unspecified
* DT_HIPROC              0x7fffffff unspecified unspecified unspecified
* "*" Signifies an entry that is at level 2.
* </pre>
* <p>
* Read more about dynamic sections at <a
href="https://refspecs.linuxbase.org/elf/gabi4+/ch5.dynamic.html#dynamic_section">
Dynamic Section</a>.
*/
public class ElfDynamicSection extends ElfSection {

    /**
     * An entry with a DT_NULL tag marks the end of the _DYNAMIC array.
     */
    public static final int DT_NULL = 0;
    /**
     * This element holds the string table offset of a null-terminated string,
giving the
     * name of a needed library. The offset is an index into the table recorded in
the
     * {@link #DT_STRTAB} code.
     * <p>
     * See <a
href="https://refspecs.linuxbase.org/elf/gabi4+/ch5.dynamic.html#shobj_dependencies">Shared Object Dependencies</a> for more information about these names.
     * <p>
     * The dynamic array may contain multiple entries with this type.

```

```

* <p>
* These entries' relative order is significant, though their relation to
entries of other types is not.
*/
public static final int DT_NEEDED = 1;
public static final int DT_PLTRELSZ = 2;
public static final int DT_PLTGOT = 3;
public static final int DT_HASH = 4;
/**
* DT_STRTAB entry holds the address, not offset, of the dynamic string table.
*/
public static final int DT_STRTAB = 5;
public static final int DT_SYMTAB = 6;
public static final int DT_RELA = 7;
public static final int DT_RELASZ = 8;
public static final int DT_RELAENT = 9;
/**
* The size in bytes of the {@link #DT_STRTAB} string table.
*/
public static final int DT_STRSZ = 10;
public static final int DT_SYMENT = 11;
public static final int DT_INIT = 12;
public static final int DT_FINI = 13;
public static final int DT_SONAME = 14;
public static final int DT_RPATH = 15;
public static final int DT_SYMBOLIC = 16;
public static final int DT_REL = 17;
public static final int DT_RELSZ = 18;
public static final int DT_RELENT = 19;
public static final int DT_PLTREL = 20;
public static final int DT_DEBUG = 21;
public static final int DT_TEXTREL = 22;
public static final int DT_JMPREL = 23;
public static final int DT_BIND_NOW = 24;
public static final int DT_INIT_ARRAY = 25;
public static final int DT_FINI_ARRAY = 26;
public static final int DT_INIT_ARRAYSZ = 27;
public static final int DT_FINI_ARRAYSZ = 28;
public static final int DT_RUNPATH = 29;
public static final int DT_FLAGS = 30;
public static final int DT_PREINIT_ARRAY = 32;
public static final int DT_GNU_HASH = 0x6ffffef5;
public static final int DT_FLAGS_1 = 0x6ffffffb;
public static final int DT_VERDEF = 0x6ffffffc; /* Address of version
definition */
public static final int DT_VERDEFNUM = 0x6ffffffd; /* Number of version
definitions */
public static final int DT_VERNEEDED = 0x6ffffffe;
public static final int DT_VERNEEDNUM = 0x6fffffff;

public static final int DF_ORIGIN = 0x1;
public static final int DF_SYMBOLIC = 0x2;
public static final int DF_TEXTREL = 0x4;
public static final int DF_BIND_NOW = 0x8;

/**
* Set RTLD_NOW for this object.
*/
public static final int DF_1_NOW = 0x00000001;

```

```

/**
 * Set RTLD_GLOBAL for this object.
 */
public static final int DF_1_GLOBAL = 0x00000002;
/**
 * Set RTLD_GROUP for this object.
 */
public static final int DF_1_GROUP = 0x00000004;
/**
 * Set RTLD_NODELETE for this object.
 */
public static final int DF_1_NODELETE = 0x00000008;
public static final int DF_1_LOADFLTR = 0x00000010;
public static final int DF_1_INITFIRST = 0x00000020;
/**
 * Object can not be used with dlopen(3)
 */
public static final int DF_1_NOOPEN = 0x00000040;
public static final int DF_1_ORIGIN = 0x00000080;
public static final int DF_1_DIRECT = 0x00000100;
public static final int DF_1_TRANS = 0x00000200;
public static final int DF_1_INTERPOSE = 0x00000400;
public static final int DF_1_NODEFLIB = 0x00000800;
/**
 * Object cannot be dumped with dldump(3)
 */
public static final int DF_1_NODUMP = 0x00001000;
public static final int DF_1_CONFALT = 0x00002000;
public static final int DF_1_ENDFILTEE = 0x00004000;
public static final int DF_1_DISPRELDNE = 0x00008000;
public static final int DF_1_DISPRELPND = 0x00010000;
public static final int DF_1_NODIRECT = 0x00020000;
public static final int DF_1_IGNMULDEF = 0x00040000;
public static final int DF_1_NOKSYMS = 0x00080000;
public static final int DF_1_NOHDR = 0x00100000;
public static final int DF_1_EDITED = 0x00200000;
public static final int DF_1_NORELOC = 0x00400000;
public static final int DF_1_SYMTNPOSE = 0x00800000;
public static final int DF_1_GLOBAUDIT = 0x01000000;
public static final int DF_1_SINGLETON = 0x02000000;
public static final int DF_1_STUB = 0x04000000;
public static final int DF_1_PIE = 0x08000000;

/**
 * For the {@link #DT_STRTAB}. Mandatory.
 */
public long dt_strtab_offset;

/**
 * For the {@link #DT_STRSZ}. Mandatory.
 */
public int dt_strtab_size;

private MemoizedObject<ElfStringTable> dtStringTable;
public final List<ElfDynamicStructure> entries = new ArrayList<>();

/**
 * An entry in the {@link #entries} of a {@link ElfDynamicSection}.
 * <p>

```


** In the elf.h header file this represents either of the following structures:*

```
*
* <pre>
* typedef struct {
*     Elf32_Sword d_tag;
*     union {
*         Elf32_Word      d_val;
*         Elf32_Addr      d_ptr;
*         Elf32_Off       d_off;
*     } d_un;
* } Elf32_Dyn;
*
* typedef struct {
*     Elf64_Xword d_tag;
*     union {
*         Elf64_Xword d_val;
*         Elf64_Addr d_ptr;
*     } d_un;
* } Elf64_Dyn;
* </pre>
*/
public static class ElfDynamicStructure {
    public ElfDynamicStructure(long d_tag, long d_val_or_ptr) {
        this.tag = d_tag;
        this.d_val_or_ptr = d_val_or_ptr;
    }

    /**
     * A tag value whose value defines how to interpret {@link #d_val_or_ptr}.
     * <p>
     * One of the DT_* constants in {@link ElfDynamicSection}.
     */
    public final long tag;

    /**
     * A field whose value is to be interpreted as specified by the {@link
#tag}.
     */
    public final long d_val_or_ptr;

    @Override
    public int hashCode() {
        final int prime = 31;
        int result = 1;
        result = prime * result + (int) (tag ^ (tag >>> 32));
        result = prime * result + (int) (d_val_or_ptr ^ (d_val_or_ptr >>>
32));
        return result;
    }

    @Override
    public boolean equals(Object obj) {
        if (this == obj) return true;
        if (obj == null) return false;
        if (getClass() != obj.getClass()) return false;
        ElfDynamicStructure other = (ElfDynamicStructure) obj;
        if (tag != other.tag) return false;
        return d_val_or_ptr == other.d_val_or_ptr;
    }
}
```

```

        @Override
        public String toString() {
            return "ElfDynamicSectionEntry{tag=" + tag + ", d_val_or_ptr=" +
d_val_or_ptr + "}";
        }
    }

    public ElfDynamicSection(final ElfParser parser, ElfSectionHeader header) {
        super(parser, header);

        parser.seek(header.section_offset);
        int numEntries = (int) (header.size / 8);

        // Except for the DT_NULL element at the end of the array, and the
        // relative order of DT_NEEDED elements, entries
        // may appear in any order. So important to use lazy evaluation to only
        // evaluating e.g. DT_STRTAB after the
        // necessary DT_STRSZ is read.
        loop:
        for (int i = 0; i < numEntries; i++) {
            long d_tag = parser.readIntOrLong();
            final long d_val_or_ptr = parser.readIntOrLong();
            entries.add(new ElfDynamicStructure(d_tag, d_val_or_ptr));
            switch ((int) d_tag) {
                case DT_NULL:
                    // A DT_NULL element ends the array (may be following DT_NULL
                    // values, but no need to look at them).
                    break loop;
                case DT_STRTAB: {
                    dtStringTable = new MemoizedObject<ElfStringTable>() {
                        @Override
                        protected ElfStringTable computeValue() throws
ElfException {
                            long fileOffsetForStringTable =
parser.virtualMemoryAddrToFileOffset(d_val_or_ptr);
                            return new ElfStringTable(parser,
fileOffsetForStringTable, dt_strtab_size, null); // FIXME: null header
                        }
                    };
                    dt_strtab_offset = d_val_or_ptr;
                }
                break;
                case DT_STRSZ:
                    if (d_val_or_ptr > Integer.MAX_VALUE) throw new
ElfException("Too large DT_STRSZ: " + d_val_or_ptr);
                    dt_strtab_size = (int) d_val_or_ptr;
                    break;
            }
        }
    }

    private ElfDynamicStructure firstEntryWithTag(long desiredTag) {
        for (ElfDynamicStructure entry : this.entries) {
            if (entry.tag == desiredTag) return entry;
        }
        return null;
    }
}

```

```

    public List<String> getNeededLibraries() throws ElfException {
        ElfStringTable stringTable = dtStringTable.getValue();
        List<String> result = new ArrayList<>();
        for (ElfDynamicStructure entry : this.entries) {
            if (entry.tag == DT_NEEDED) result.add(stringTable.get((int)
entry.d_val_or_ptr));
        }
        return result;
    }

    public String getRunPath() {
        ElfDynamicStructure runPathEntry = firstEntryWithTag(DT_RUNPATH);
        return runPathEntry == null ? null : dtStringTable.getValue().get((int)
runPathEntry.d_val_or_ptr);
    }

    public long getFlags() {
        ElfDynamicStructure flagsEntry = firstEntryWithTag(DT_FLAGS);
        return flagsEntry == null ? 0 : flagsEntry.d_val_or_ptr;
    }

    public long getFlags1() {
        ElfDynamicStructure flagsEntry = firstEntryWithTag(DT_FLAGS_1);
        return flagsEntry == null ? 0 : flagsEntry.d_val_or_ptr;
    }

    @Override
    public String toString() {
        return "ElfDynamicStructure{entries=" + this.entries + "}";
    }
}

```

src/net/fornwall/jelf/ElfException.java

```
package net.fornwall.jelf;
```

```

/**
 * Generic exception class for all exceptions which occur in this package. Since
 * there is no mechanism built into this library for recovering from errors, the
 * best clients can do is display the error string.
 */

```

```

public class ElfException extends RuntimeException {

    private static final long serialVersionUID = 1L;

    public ElfException(String message) {
        super(message);
    }

    public ElfException(Throwable cause) {
        super(cause);
    }

    public ElfException(String message, Throwable cause) {
        super(message, cause);
    }
}

```

src/net/fornwall/jelf/ElfFile.java

```
package net.fornwall.jelf;

import java.io.ByteArrayInputStream;
import java.io.ByteArrayOutputStream;
import java.io.File;
import java.io.FileInputStream;
import java.io.IOException;
import java.io.InputStream;
import java.nio.MappedByteBuffer;
import java.util.ArrayList;
import java.util.Collections;
import java.util.List;

/**
 * An ELF (Executable and Linkable Format) file that can be a relocatable,
 * executable, shared or core file.
 * <p>
 * Use one of the following methods to parse input to get an instance of this
 * class:
 * <ul>
 * <li>{@link #from(File)}</li>
 * <li>{@link #from(byte[])}</li>
 * <li>{@link #from(InputStream)}</li>
 * <li>{@link #from(MappedByteBuffer)}</li>
 * </ul>
 * <p>
 * Resources about ELF files:
 * <ul>
 * <li>http://man7.org/linux/man-pages/man5/elf.5.html</li>
 * <li>http://en.wikipedia.org/wiki/Executable\_and\_Linkable\_Format</li>
 * <li>http://www.ibm.com/developerworks/library/l-dynamic-libraries/</li>
 * <li>http://downloads.openwatcom.org/ftp/devel/docs/elf-64-gen.pdf</li>
 * </ul>
 */
public final class ElfFile {

    /**
     * Relocatable file type. A possible value of {@link #e_type}.
     */
    public static final int ET_REL = 1;

    /**
     * Executable file type. A possible value of {@link #e_type}.
     */
    public static final int ET_EXEC = 2;

    /**
     * Shared object file type. A possible value of {@link #e_type}.
     */
    public static final int ET_DYN = 3;

    /**
     * Core file file type. A possible value of {@link #e_type}.
     */
    public static final int ET_CORE = 4;

    /**
     * 32-bit objects.
     */
}
```

```

public static final byte CLASS_32 = 1;
/**
 * 64-bit objects.
 */
public static final byte CLASS_64 = 2;

/**
 * LSB data encoding.
 */
public static final byte DATA_LSB = 1;
/**
 * MSB data encoding.
 */
public static final byte DATA_MSB = 2;

/**
 * No architecture type.
 */
public static final int ARCH_NONE = 0;
/**
 * AT&T architecture type.
 */
public static final int ARCH_ATT = 1;
/**
 * SPARC architecture type.
 */
public static final int ARCH_SPARC = 2;
/**
 * Intel 386 architecture type.
 */
public static final int ARCH_i386 = 3;
/**
 * Motorola 68000 architecture type.
 */
public static final int ARCH_68k = 4;
/**
 * Motorola 88000 architecture type.
 */
public static final int ARCH_88k = 5;
/**
 * Intel 860 architecture type.
 */
public static final int ARCH_i860 = 7;
/**
 * MIPS architecture type.
 */
public static final int ARCH_MIPS = 8;
public static final int ARCH_ARM = 0x28;
public static final int ARCH_X86_64 = 0x3E;
public static final int ARCH_AARCH64 = 0xB7;

/**
 * Identifies the object file type. One of the ET_* constants in the class.
 */
public final short e_type; // Elf32_Half
/**
 * Byte identifying the size of objects, either {@link #CLASS_32} or {@link
 * #CLASS_64} .
 */

```

```

    public final byte objectSize;

    /**
     * Returns a byte identifying the data encoding of the processor specific
     data. This byte will be either
     * DATA_INVALID, DATA_LSB or DATA_MSB.
     */
    public final byte encoding;

    public final byte elfVersion;
    public final byte abi;
    public final byte abiVersion;

    /**
     * The required architecture. One of the ARCH_* constants in the class.
     */
    public final short arch; // Elf32_Half

    /**
     * Version
     */
    public final int version; // Elf32_Word

    /**
     * Virtual address to which the system first transfers control. If there is no
     entry point for the file the value is
     * 0.
     */
    public final long entry_point; // Elf32_Addr

    /**
     * e_phoff. Program header table offset in bytes. If there is no program
     header table the value is 0.
     */
    public final long ph_offset; // Elf32_Off

    /**
     * e_shoff. Section header table offset in bytes. If there is no section
     header table the value is 0.
     */
    public final long sh_offset; // Elf32_Off

    /**
     * e_flags. Processor specific flags.
     */
    public final int flags; // Elf32_Word

    /**
     * e_ehsize. ELF header size in bytes.
     */
    public final short eh_size; // Elf32_Half

    /**
     * e_phentsize. Size of one entry in the file's program header table in bytes.
     All entries are the same size.
     */
    public final short ph_entry_size; // Elf32_Half

    /**
     * e_phnum. Number of {@link ElfSegment} entries in the program header table,
     0 if no entries.
     */
    public final short num_ph; // Elf32_Half

    /**
     * e_shentsize. Section header entry size in bytes - all entries are the same
     size.
     */

```

```

public final short sh_entry_size; // Elf32_Half
/**
 * e_shnum. Number of entries in the section header table, 0 if no entries.
 */
public final short num_sh; // Elf32_Half

/**
 * Elf{32,64}_Ehdr#e_shstrndx. Index into the section header table associated
with the section name string table.
 * SH_UNDEF if there is no section name string table.
 */
private short sh_string_ndx; // Elf32_Half

/**
 * MemoizedObject array of section headers associated with this ELF file.
 */
private MemoizedObject<ElfSection>[] sections;
/**
 * MemoizedObject array of program headers associated with this ELF file.
 */
private MemoizedObject<ElfSegment>[] programHeaders;

/**
 * Used to cache symbol table lookup.
 */
private ElfSymbolTableSection symbolTableSection;
/**
 * Used to cache dynamic symbol table lookup.
 */
private ElfSymbolTableSection dynamicSymbolTableSection;

private ElfDynamicSection dynamicSection;

/**
 * Returns the section header at the specified index. The section header at
index 0 is defined as being a undefined
 * section.
 */
public ElfSection getSection(int index) throws ElfException {
    return sections[index].getValue();
}

public List<ElfSection> sectionsOfType(int sectionType) throws ElfException {
    if (num_sh < 2) return Collections.emptyList();
    List<ElfSection> result = new ArrayList<>();
    for (int i = 1; i < num_sh; i++) {
        ElfSection section = getSection(i);
        if (section.header.type == sectionType) {
            result.add(section);
        }
    }
    return result;
}

/**
 * Returns the section header string table associated with this ELF file.
 */
public ElfStringTable getSectionNameStringTable() throws ElfException {

```

```

        return (ElfStringTable) getSection(sh_string_ndx);
    }

    /**
     * Returns the string table associated with this ELF file.
     */
    public ElfStringTable getStringTable() throws ElfException {
        return findStringTableWithName(ElfSectionHeader.NAME_STRTAB);
    }

    /**
     * Returns the dynamic symbol table associated with this ELF file, or null if
     one does not exist.
     */
    public ElfStringTable getDynamicStringTable() throws ElfException {
        return findStringTableWithName(ElfSectionHeader.NAME_DYNSTR);
    }

    private ElfStringTable findStringTableWithName(String tableName) throws
    ElfException {
        // Loop through the section header and look for a section
        // header with the name "tableName". We can ignore entry 0
        // since it is defined as being undefined.
        return (ElfStringTable) firstSectionByName(tableName);
    }

    /**
     * The {@link ElfSectionHeader#SHT_SYMTAB} section (of which there may be only
     one), if any.
     */
    public ElfSymbolTableSection getSymbolTableSection() throws ElfException {
        return (symbolTableSection != null) ? symbolTableSection :
        (symbolTableSection = (ElfSymbolTableSection)
        firstSectionByType(ElfSectionHeader.SHT_SYMTAB));
    }

    /**
     * The {@link ElfSectionHeader#SHT_DYNSYM} section (of which there may be only
     one), if any.
     */
    public ElfSymbolTableSection getDynamicSymbolTableSection() throws
    ElfException {
        return (dynamicSymbolTableSection != null) ? dynamicSymbolTableSection :
        (dynamicSymbolTableSection = (ElfSymbolTableSection)
        firstSectionByType(ElfSectionHeader.SHT_DYNSYM));
    }

    /**
     * The {@link ElfSectionHeader#SHT_DYNAMIC} section (of which there may be
     only one). Named ".dynamic".
     */
    public ElfDynamicSection getDynamicSection() {
        return (dynamicSection != null) ? dynamicSection : (dynamicSection =
        (ElfDynamicSection) firstSectionByType(ElfSectionHeader.SHT_DYNAMIC));
    }

    public ElfSection firstSectionByType(int type) throws ElfException {
        for (int i = 1; i < num_sh; i++) {
            ElfSection sh = getSection(i);

```



```

        if (sh.header.type == type) return sh;
    }
    return null;
}

public <T extends ElfSection> T firstSectionByType(Class<T> type) throws
ElfException {
    for (int i = 1; i < num_sh; i++) {
        ElfSection sh = getSection(i);
        if (type.isInstance(sh)) return (T) sh;
    }
    return null;
}

public ElfSection firstSectionByName(String sectionName) throws ElfException {
    for (int i = 1; i < num_sh; i++) {
        ElfSection sh = getSection(i);
        if (sectionName.equals(sh.header.getName())) return sh;
    }
    return null;
}

/**
 * Returns the elf symbol with the specified name or null if one is not found.
 */
public ElfSymbol getELFSymbol(String symbolName) throws ElfException,
IOException {
    if (symbolName == null) return null;

    // Check dynamic symbol table for symbol name.
    ElfSymbolTableSection sh = getDynamicSymbolTableSection();
    if (sh != null) {
        int numSymbols = sh.symbols.length;
        for (int i = 0; i < Math.ceil(numSymbols / 2); i++) {
            ElfSymbol symbol = sh.symbols[i];
            if (symbolName.equals(symbol.getName())) {
                return symbol;
            } else if (symbolName.equals((symbol = sh.symbols[numSymbols - 1 -
i]).getName())) {
                return symbol;
            }
        }
    }

    // Check symbol table for symbol name.
    sh = getSymbolTableSection();
    if (sh != null) {
        int numSymbols = sh.symbols.length;
        for (int i = 0; i < Math.ceil(numSymbols / 2); i++) {
            ElfSymbol symbol = sh.symbols[i];
            if (symbolName.equals(symbol.getName())) {
                return symbol;
            } else if (symbolName.equals((symbol = sh.symbols[numSymbols - 1 -
i]).getName())) {
                return symbol;
            }
        }
    }
    return null;
}

```

```

    }

    /**
     * Returns the elf symbol with the specified address or null if one is not
     found. 'address' is relative to base of
     * shared object for .so's.
     */
    public ElfSymbol getELFSymbol(long address) throws ElfException {
        // Check dynamic symbol table for address.
        ElfSymbol symbol;
        long value;

        ElfSymbolTableSection sh = getDynamicSymbolTableSection();
        if (sh != null) {
            int numSymbols = sh.symbols.length;
            for (int i = 0; i < numSymbols; i++) {
                symbol = sh.symbols[i];
                value = symbol.st_value;
                if (address >= value && address < value + symbol.st_size) return
symbol;
            }
        }

        // Check symbol table for symbol name.
        sh = getSymbolTableSection();
        if (sh != null) {
            int numSymbols = sh.symbols.length;
            for (int i = 0; i < numSymbols; i++) {
                symbol = sh.symbols[i];
                value = symbol.st_value;
                if (address >= value && address < value + symbol.st_size) return
symbol;
            }
        }
        return null;
    }

    public ElfSegment getProgramHeader(int index) {
        return programHeaders[index].getValue();
    }

    public static ElfFile from(InputStream in) throws IOException {
        ByteArrayOutputStream baos = new ByteArrayOutputStream();
        int totalRead = 0;
        byte[] buffer = new byte[8096];
        boolean firstRead = true;
        while (true) {
            int readNow = in.read(buffer, totalRead, buffer.length - totalRead);
            if (readNow == -1) {
                return from(baos.toByteArray());
            } else {
                if (firstRead) {
                    // Abort early.
                    if (readNow < 4) {
                        throw new ElfException("Bad first read");
                    } else {
                        if (!(0x7f == buffer[0] && 'E' == buffer[1] && 'L' ==
buffer[2] && 'F' == buffer[3]))
                            throw new ElfException("Bad magic number for file");
                    }
                }
            }
        }
    }

```

```

        }
        firstRead = false;
    }
    baos.write(buffer, 0, readNow);
}
}

public static ElfFile from(File file) throws ElfException, IOException {
    byte[] buffer = new byte[(int) file.length()];
    try (FileInputStream in = new FileInputStream(file)) {
        int totalRead = 0;
        while (totalRead < buffer.length) {
            int readNow = in.read(buffer, totalRead, buffer.length -
totalRead);
            if (readNow == -1) {
                throw new ElfException("Premature end of file");
            } else {
                totalRead += readNow;
            }
        }
        return from(buffer);
    }
}

public static ElfFile from(byte[] buffer) throws ElfException, IOException {
    return new ElfFile(new BackingFile(new ByteArrayInputStream(buffer)));
}

public static ElfFile from(MappedByteBuffer mappedByteBuffer) throws
ElfException, IOException {
    return new ElfFile(new BackingFile(mappedByteBuffer));
}

public final ElfParser parser;

private ElfFile(BackingFile backingFile) throws ElfException, IOException {
    parser = new ElfParser(this, backingFile);

    byte[] ident = new byte[16];
    int bytesRead = parser.read(ident);
    if (bytesRead != ident.length)
        throw new ElfException("Error reading elf header (read " + bytesRead +
"bytes - expected to read " + ident.length + "bytes)");

    if (!(0x7f == ident[0] && 'E' == ident[1] && 'L' == ident[2] && 'F' ==
ident[3]))
        throw new ElfException("Bad magic number for file");

    objectSize = ident[4];
    if (!(objectSize == CLASS_32 || objectSize == CLASS_64))
        throw new ElfException("Invalid object size class: " + objectSize);
    encoding = ident[5];
    if (!(encoding == DATA_LSB || encoding == DATA_MSB)) throw new
ElfException("Invalid encoding: " + encoding);
    elfVersion = ident[6];
    if (elfVersion != 1) throw new ElfException("Invalid elf version: " +
elfVersion);
    abi = ident[7]; // EI_OSABI, target operating system ABI

```

```

abiVersion = ident[8]; // EI_ABIVERSION, ABI version. Linux kernel (after
at least 2.6) has no definition of it.
// ident[9-15] // EI_PAD, currently unused.

```

```

e_type = parser.readShort();
arch = parser.readShort();
version = parser.readInt();
entry_point = parser.readIntOrLong();
ph_offset = parser.readIntOrLong();
sh_offset = parser.readIntOrLong();
flags = parser.readInt();
eh_size = parser.readShort();
ph_entry_size = parser.readShort();
num_ph = parser.readShort();
sh_entry_size = parser.readShort();
num_sh = parser.readShort();
if (num_sh == 0) {
    throw new ElfException("e_shnum is SHN_UNDEF(0), which is not
supported yet"
        + " (the actual number of section header table entries is
contained in the sh_size field of the section header at index 0)");
}
sh_string_ndx = parser.readShort();
if (sh_string_ndx == /* SHN_XINDEX= */0xffff) {
    throw new ElfException("e_shstrndx is SHN_XINDEX(0xffff), which is not
supported yet"
        + " (the actual index of the section name string table section
is contained in the sh_link field of the section header at index 0)");
}

sections = MemoizedObject.uncheckedArray(num_sh);
for (int i = 0; i < num_sh; i++) {
    final long sectionHeaderOffset = sh_offset + (i * sh_entry_size);
    sections[i] = new MemoizedObject<>() {
        @Override
        public ElfSection computeValue() throws ElfException {
            ElfSectionHeader elfSectionHeader = new
ElfSectionHeader(parser, sectionHeaderOffset);
            switch (elfSectionHeader.type) {
                case ElfSectionHeader.SHT_DYNAMIC:
                    return new ElfDynamicSection(parser,
elfSectionHeader);

                case ElfSectionHeader.SHT_SYMTAB:
                case ElfSectionHeader.SHT_DYNSYM:
                    return new ElfSymbolTableSection(parser,
elfSectionHeader);

                case ElfSectionHeader.SHT_STRTAB:
                    return new ElfStringTable(parser,
elfSectionHeader.section_offset, (int) elfSectionHeader.size, elfSectionHeader);
                case ElfSectionHeader.SHT_HASH:
                    return new ElfHashTable(parser, elfSectionHeader);
                case ElfSectionHeader.SHT_NOTE:
                    return new ElfNoteSection(parser, elfSectionHeader);
                case ElfSectionHeader.SHT_RELA:
                    return new ElfRelocationSection(parser,
elfSectionHeader);

                case ElfSectionHeader.SHT_GNU_HASH:
                    return new ElfGnuHashTable(parser, elfSectionHeader);
                default:

```

```

        return new ElfSection(parser, elfSectionHeader);
    }
}

};

}

programHeaders = MemoizedObject.uncheckedArray(num_ph);
for (int i = 0; i < num_ph; i++) {
    final long programHeaderOffset = ph_offset + (i * ph_entry_size);
    programHeaders[i] = new MemoizedObject<ElfSegment>() {
        @Override
        public ElfSegment computeValue() {
            return new ElfSegment(parser, programHeaderOffset);
        }
    };
}

}

/**
 * The interpreter specified by the {@link ElfSegment#PT_INTERP} program
 * header, if any.
 */
public String getInterpreter() throws IOException {
    for (MemoizedObject<ElfSegment> programHeader : programHeaders) {
        ElfSegment ph = programHeader.getValue();
        if (ph.type == ElfSegment.PT_INTERP) return ph.getIntepreter();
    }
    return null;
}

}

```

src/net/fornwall/jelf/ElfGnuHashTable.java

```

package net.fornwall.jelf;

/**
 * An ELF section containing a hash table for lookup of dynamic symbols.
 *
 * Has the section type {@link ElfSectionHeader#SHT_GNU_HASH}.
 *
 * Replaces {@link ElfHashTable} on almost all modern Linux systems.
 *
 * See https://flapenguin.me/2017/05/10/elf-lookup-dt-gnu-hash/
 */
public class ElfGnuHashTable extends ElfSection {

    private final ElfParser parser;
    private final int ELFCLASS_BITS;
    // The number of .dynsym symbols skipped.
    int symbolOffset;
    int bloomShift;
    long[] bloomFilter;
    int[] buckets;
    int[] chain;

    ElfGnuHashTable(ElfParser parser, ElfSectionHeader header) {
        super(parser, header);
        this.parser = parser;
    }
}

```

```

    ELFCLASS_BITS = parser.elfFile.objectSize == ElfFile.CLASS_32 ? 32 : 64;

    parser.seek(header.section_offset);
    int numberOfBuckets = parser.readInt();
    symbolOffset = parser.readInt();
    int bloomSize = parser.readInt();
    bloomShift = parser.readInt();
    bloomFilter = new long[bloomSize];
    buckets = new int[numberOfBuckets];

    for (int i = 0; i < bloomSize; i++) {
        bloomFilter[i] = parser.readIntOrLong();
    }
    for (int i = 0; i < numberOfBuckets; i++) {
        buckets[i] = parser.readInt();
    }
    // The chain is initialized on first use in lookupSymbol() due to it
    requiring .dynsym size.
}

ElfSymbol lookupSymbol(String symbolName, ElfSymbolTableSection symbolTable) {
    if (chain == null) {
        int chainSize = ((ElfSymbolTableSection)
parser.elfFile.firstSectionByType(ElfSectionHeader.SHT_DYNSYM)).symbols.length -
symbolOffset;
        chain = new int[chainSize];
        parser.seek(header.section_offset + 4*4 +
bloomFilter.length*(ELFCLASS_BITS/8) + buckets.length * 4);
        for (int i = 0; i < chainSize; i++) {
            chain[i] = parser.readInt();
        }
    }

    final int nameHash = gnuHash(symbolName);

    long word =
bloomFilter[(Integer remainderUnsigned(Integer.divideUnsigned(nameHash,
ELFCLASS_BITS), bloomFilter.length))];
    long mask = 1L << (long) (Integer remainderUnsigned(nameHash,
ELFCLASS_BITS))
        | 1L << (long) (Integer remainderUnsigned((nameHash >>>
bloomShift), ELFCLASS_BITS));

    if ((word & mask) != mask) {
        // If at least one bit is not set, a symbol is surely missing.
        return null;
    }

    int symix = buckets[Integer remainderUnsigned(nameHash, buckets.length)];
    if (symix < symbolOffset) {
        return null;
    }

    while (true) {
        int hash = chain[symix - symbolOffset];

        if (((long) nameHash | 1L) == ((long) hash | 1L)) {
            // The chain contains contiguous sequences of hashes for symbols

```

```

hashing to the same index,
    // with the lowest bit discarded (used to signal end of chain).
    ElfSymbol symbol = symbolTable.symbols[symix];
    if (symbolName.equals(symbol.getName())) return symbol;
}
ElfSymbol symbol = symbolTable.symbols[symix];

if ((hash & 1) != 0) {
    // Chain ends with an element with the lowest bit set to 1.
    break;
}

symix++;
}

return null;
}

static int gnuHash(String name) {
    int h = 5381;
    int nameLength = name.length();
    for (int i = 0; i < nameLength; i++) {
        char c = name.charAt(i);
        h = (h << 5) + h + c;
    }
    return h;
}
}

```

src/net/fornwall/jelf/ElfHashTable.java

```

package net.fornwall.jelf;

/**
 * An ELF section containing a hash table for lookup of dynamic symbols.
 *
 * Note that this has been replaced with {@link ElfGnuHashTable} on modern Linux
 * systems.
 *
 * See https://flapenguin.me/2017/04/24/elf-lookup-dt-hash/
 */
public class ElfHashTable extends ElfSection {

    private final int[] buckets;
    private final int[] chain;

    ElfHashTable(ElfParser parser, ElfSectionHeader header) {
        super(parser, header);

        parser.seek(header.section_offset);

        int num_buckets = parser.readInt();
        int num_chains = parser.readInt();

        buckets = new int[num_buckets];
        for (int i = 0; i < num_buckets; i++) {
            buckets[i] = parser.readInt();
        }
    }
}

```

```

        chain = new int[num_chains];
        for (int i = 0; i < num_chains; i++) {
            chain[i] = parser.readInt();
        }

        // Make sure that the amount of bytes we were supposed to read
        // was what we actually read.
        int actual = num_buckets * 4 + num_chains * 4 + 8;
        if (header.size != actual) {
            throw new ElfException("Error reading string table (read " + actual +
"bytes, expected to read " + header.size + "bytes).");
        }
    }

    public ElfSymbol lookupSymbol(String name, ElfSymbolTableSection symbolTable)
    {
        long hashValue = elfHash(name);
        int index = buckets[(int) (hashValue % buckets.length)];
        while (true) {
            if (index == 0) return null;
            ElfSymbol symbol = symbolTable.symbols[index];
            if (name.equals(symbol.getName())) return symbol;
            index = chain[index];
        }
    }

    static long elfHash(String name) {
        long hash = 0;
        int nameLength = name.length();
        for (int i = 0; i < nameLength; i++) {
            hash = (hash << 4) + name.charAt(i);
            long x = hash & 0xF0000000L;
            if (x != 0) hash ^= (x >> 24);
            hash &= ~x;
        }
        return hash;
    }
}

```

src/net/fornwall/jelf/ElfNoteSection.java

```

package net.fornwall.jelf;

import java.io.IOException;

class ElfNoteSection extends ElfSection {

    /**
     * A possible value of the {@link #type} where the description should contain
     * {@link GnuAbiDescriptor}.
     */
    public static final int NT_GNU_ABI_TAG = 1;

    /**
     * A possible value of the {@link #type} for a note containing synthetic hwcap
     * information.
     *
     * The descriptor begins with two words:
     * word 0: number of entries
     */
}

```



```

    *    word 1: bitmask of enabled entries
    *    Then follow variable-length entries, one byte followed by a '\0'-
    *    terminated hwcap name string. The byte gives the bit
    *    number to test if enabled, (1U << bit) & bitmask.
    */
    public static final int NT_GNU_HWCAP = 2;
    /**
     * A possible value of the {@link #type} for a note containing build ID bits
     * as generated by "ld --build-id".
     *
     * The descriptor consists of any nonzero number of bytes.
     */
    public static final int NT_GNU_BUILD_ID = 3;

    /**
     * A possible value of the {@link #type} for a note containing a version
     * string generated by GNU gold.
     */
    public static final int NT_GNU_GOLD_VERSION = 4;

    /**
     * The descriptor content of a Link {@link #NT_GNU_ABI_TAG} type note.
     *
     * Accessible in {@link #descriptorAsGnuAbi()}.
     */
    public final static class GnuAbiDescriptor {

        /** A possible value of {@link #operatingSystem}. */
        public static final int ELF_NOTE_OS_LINUX = 0;
        /** A possible value of {@link #operatingSystem}. */
        public static final int ELF_NOTE_OS_GNU = 1;
        /** A possible value of {@link #operatingSystem}. */
        public static final int ELF_NOTE_OS_SOLARIS2 = 2;
        /** A possible value of {@link #operatingSystem}. */
        public static final int ELF_NOTE_OS_FREEBSD = 3;

        /** One of the ELF_NOTE_OS_* constants in this class. */
        public final int operatingSystem;
        /** Major version of the required ABI. */
        public final int majorVersion;
        /** Minor version of the required ABI. */
        public final int minorVersion;
        /** Subminor version of the required ABI. */
        public final int subminorVersion;

        public GnuAbiDescriptor(int operatingSystem, int majorVersion, int
minorVersion, int subminorVersion) {
            this.operatingSystem = operatingSystem;
            this.majorVersion = majorVersion;
            this.minorVersion = minorVersion;
            this.subminorVersion = subminorVersion;
        }
    }

    public final /* uint32_t */ int nameSize;
    public final /* uint32_t */ int descriptorSize;
    public final /* uint32_t */ int type;
    private String name;
    private byte[] descriptorBytes;

```

```

private final GnuAbiDescriptor gnuAbiDescriptor;

ElfNoteSection(ElfParser parser, ElfSectionHeader header) throws ElfException
{
    super(parser, header);

    parser.seek(header.section_offset);
    nameSize = parser.readInt();
    descriptorSize = parser.readInt();
    type = parser.readInt();
    byte[] nameBytes = new byte[nameSize];
    descriptorBytes = new byte[descriptorSize];
    int bytesRead = parser.read(nameBytes);
    if (bytesRead != nameSize) {
        throw new ElfException("Error reading note name (read=" + bytesRead +
            ", expected=" + nameSize + ")");
    }
    parser.skip(bytesRead % 4);

    switch (type) {
        case NT_GNU_ABI_TAG:
            gnuAbiDescriptor = new GnuAbiDescriptor(parser.readInt(),
                parser.readInt(), parser.readInt(), parser.readInt());
            break;
        default:
            gnuAbiDescriptor = null;
    }

    bytesRead = parser.read(descriptorBytes);
    if (bytesRead != descriptorSize) {
        throw new ElfException("Error reading note name (read=" + bytesRead +
            ", expected=" + descriptorSize + ")");
    }

    name = new String(nameBytes, 0, nameSize-1); // unnecessary trailing 0
}

String getName() {
    return name;
}

byte[] descriptorBytes() {
    return descriptorBytes;
}

public String descriptorAsString() {
    return new String(descriptorBytes);
}

public GnuAbiDescriptor descriptorAsGnuAbi() {
    return gnuAbiDescriptor;
}
}

```

src/net/fornwall/jelf/ElfParser.java

```
package net.fornwall.jelf;
```

```

/**
 * Package internal class used for parsing ELF files.
 */
public class ElfParser {

    final ElfFile elfFile;
    private final BackingFile backingFile;
    private long readBytes;

    ElfParser(ElfFile elfFile, BackingFile backingFile) {
        this.elfFile = elfFile;
        this.backingFile = backingFile;
    }

    public void seek(long offset) {
        readBytes = 0;
        backingFile.seek(offset);
    }

    public void skip(int bytesToSkip) {
        readBytes = 0;
        backingFile.skip(bytesToSkip);
    }

    public long getReadBytes() {
        return readBytes;
    }

    /**
     * Signed byte utility functions used for converting from big-endian (MSB) to
     * little-endian (LSB).
     */
    short byteSwap(short arg) {
        return (short) ((arg << 8) | ((arg >>> 8) & 0xFF));
    }

    int byteSwap(int arg) {
        return ((byteSwap((short) arg)) << 16) | (((byteSwap((short) (arg >>>
16)))) & 0xFFFF);
    }

    long byteSwap(long arg) {
        return (((long) byteSwap((int) arg)) << 32) | (((long) byteSwap((int)
(arg >>> 32))) & 0xFFFFFFFF));
    }

    short readUnsignedByte() {
        readBytes++;
        return backingFile.readUnsignedByte();
    }

    public short readShort() throws ElfException {
        int ch1 = readUnsignedByte();
        int ch2 = readUnsignedByte();
        short val = (short) ((ch1 << 8) + (ch2 << 0));
        if (elfFile.encoding == ElfFile.DATA_LSB) val = byteSwap(val);
        return val;
    }
}

```

```

public int readInt() throws ElfException {
    int ch1 = readUnsignedByte();
    int ch2 = readUnsignedByte();
    int ch3 = readUnsignedByte();
    int ch4 = readUnsignedByte();
    int val = ((ch1 << 24) + (ch2 << 16) + (ch3 << 8) + (ch4));

    if (elfFile.encoding == ElfFile.DATA_LSB) val = byteSwap(val);
    return val;
}

public long readLong() {
    int ch1 = readUnsignedByte();
    int ch2 = readUnsignedByte();
    int ch3 = readUnsignedByte();
    int ch4 = readUnsignedByte();
    int val1 = ((ch1 << 24) + (ch2 << 16) + (ch3 << 8) + (ch4 << 0));
    int ch5 = readUnsignedByte();
    int ch6 = readUnsignedByte();
    int ch7 = readUnsignedByte();
    int ch8 = readUnsignedByte();
    int val2 = ((ch5 << 24) + (ch6 << 16) + (ch7 << 8) + (ch8 << 0));

    long val = ((long) (val1) << 32) + (val2 & 0xFFFFFFFFL);
    if (elfFile.encoding == ElfFile.DATA_LSB) val = byteSwap(val);
    return val;
}

/**
 * Read four-byte int or eight-byte Long depending on if {@link
 * ElfFile#objectSize}.
 */
public long readIntOrLong() {
    return elfFile.objectSize == ElfFile.CLASS_32 ? readInt() : readLong();
}

/**
 * Returns a big-endian unsigned representation of the int.
 */
public long unsignedByte(int arg) {
    long val;
    if (arg >= 0) {
        val = arg;
    } else {
        val = (unsignedByte((short) (arg >>> 16)) << 16) | ((short) arg);
    }
    return val;
}

/**
 * Find the file offset from a virtual address by looking up the {@link
 * ElfSegment} segment containing the
 * address and computing the resulting file offset.
 */
long virtualMemoryAddrToFileOffset(long address) {
    for (int i = 0; i < elfFile.num_ph; i++) {
        ElfSegment ph = elfFile.getProgramHeader(i);
        if (address >= ph.virtual_address && address < (ph.virtual_address +
        ph.mem_size)) {

```

```

        long relativeOffset = address - ph.virtual_address;
        if (relativeOffset >= ph.file_size)
            throw new ElfException("Can not convert virtual memory address
" + Long.toHexString(address) + " to file offset -" + " found segment " + ph
            + " but address maps to memory outside file range");
        return ph.offset + relativeOffset;
    }
}
throw new ElfException("Cannot find segment for address " +
Long.toHexString(address));
}

public int read(byte[] data) {
    return backingFile.read(data);
}
}

```

src/net/fornwall/jelf/ElfRelocationSection.java

```

package net.fornwall.jelf;

public class ElfRelocationSection extends ElfSection {

    public ElfRelocationSection(ElfParser parser, ElfSectionHeader header) {
        super(parser, header);

        int num_entries = (int) (header.size / header.entry_size);
    }
}

```

src/net/fornwall/jelf/ElfSection.java

```

package net.fornwall.jelf;

public class ElfSection {
    public final ElfSectionHeader header;
    private final ElfParser parser;

    public ElfSection(ElfParser parser, ElfSectionHeader header) {
        this.header = header;
        this.parser = parser;
    }

    public byte[] rawSection() {
        parser.seek(header.section_offset);
        byte[] data = new byte[(int) header.size];
        parser.read(data);
        return data;
    }
}

```

src/net/fornwall/jelf/ElfSectionHeader.java

```

package net.fornwall.jelf;

import java.io.IOException;

```

```

/**

```

```

* Class corresponding to the Elf32_Shdr/Elf64_Shdr struct.
*
* <p>
* An object file's section header table lets one locate all the file's sections.
The section header table is an array
* of Elf32_Shdr or Elf64_Shdr structures. A section header table index is a
subscript into this array. The ELF header's
* {@Link ElfFile#sh_offset e_shoff member} gives the byte offset from the
beginning of the file to the section header
* table with each section header entry being {@Link ElfFile#sh_entry_size
e_shentsize} bytes big.
*
* <p>
* {@Link ElfFile#num_sh e_shnum} normally tells how many entries the section
header table contains, but if the number
* of sections is greater than or equal to SHN_LORESERVE (0xff00), e_shnum has the
value SHN_UNDEF (0) and the actual
* number of section header table entries is contained in the sh_size field of the
section header at index 0 (otherwise,
* the sh_size member of the initial entry contains 0).
*
* <p>
* Some section header table indexes are reserved in contexts where index size is
restricted, for example, the st_shndx
* member of a symbol table entry and the e_shnum and e_shstrndx members of the
ELF header. In such contexts, the
* reserved values do not represent actual sections in the object file. Also in
such contexts, an escape value indicates
* that the actual section index is to be found elsewhere, in a larger field.
*/
public class ElfSectionHeader {

    /**
     * Marks the section header as inactive; it does not have an associated
     section. Other members of the section header
     * have undefined values.
     */
    public static final int SHT_NULL = 0;
    /**
     * Section holds information defined by the program.
     */
    public static final int SHT_PROGBITS = 1;
    /**
     * The {@Link #type} value for a section containing complete symbol table
     information necessary for link editing.
     * <p>
     * See {@Link ElfSymbolTableSection}, which is the class representing sections
     of this type, for more information.
     */
    public static final int SHT_SYMTAB = 2;
    /**
     * Section holds string table information.
     */
    public static final int SHT_STRTAB = 3;
    /**
     * Section holds relocation entries with explicit addends.
     */
    public static final int SHT_RELA = 4;
    /**

```

```

    * Section holds symbol hash table.
    */
    public static final int SHT_HASH = 5;
    /**
    * Section holds information for dynamic linking. Only one per ELF file. The
    dynsym is allocable, and contains the
    * symbols needed to support runtime operation.
    */
    public static final int SHT_DYNAMIC = 6;
    /**
    * Section holds information that marks the file.
    */
    public static final int SHT_NOTE = 7;
    /**
    * Section occupies no space but resembles TYPE_PROGBITS.
    */
    public static final int SHT_NOBITS = 8;
    /**
    * Section holds relocation entries without explicit addends.
    */
    public static final int SHT_REL = 9;
    /**
    * Section is reserved but has unspecified semantics.
    */
    public static final int SHT_SHLIB = 10;
    /**
    * The {@link #type} value for a section containing a minimal set of symbols
    needed for dynamic linking at runtime.
    * <p>
    * See {@link ElfSymbolTableSection}, which is the class representing sections
    of this type, for more information.
    */
    public static final int SHT_DYNSYM = 11;
    public static final int SHT_INIT_ARRAY = 14;
    public static final int SHT_FINI_ARRAY = 15;
    public static final int SHT_PREINIT_ARRAY = 16;
    public static final int SHT_GROUP = 17;
    public static final int SHT_SYMTAB_SHNDX = 18;

    /**
    * A hash table for fast lookup of dynamic symbols.
    * <p>
    * See {@link ElfGnuHashTable}.
    */
    public static final int SHT_GNU_HASH = 0x6ffffff6;
    public static final int SHT_GNU_verdef = 0x6ffffffd;
    public static final int SHT_GNU_verneed = 0x6ffffffe;
    public static final int SHT_GNU_versym = 0x6fffffff;

    /**
    * Lower bound of the range of indexes reserved for operating system-specific
    semantics.
    */
    public static final int SHT_LOOS = 0x60000000;
    /**
    * Upper bound of the range of indexes reserved for operating system-specific
    semantics.
    */
    public static final int SHT_HIOS = 0x6fffffff;

```

```

/**
 * Lower bound of the range of indexes reserved for processor-specific
 semantics.
 */
public static final int SHT_LOPROC = 0x70000000;
/**
 * Upper bound of the range of indexes reserved for processor-specific
 semantics.
 */
public static final int SHT_HIPROC = 0x7fffffff;
/**
 * Lower bound of the range of indexes reserved for application programs.
 */
public static final int SHT_LOUSER = 0x80000000;
/**
 * Upper bound of the range of indexes reserved for application programs.
 */
public static final int SHT_HIUSER = 0xffffffff;

public static final short SHN_UNDEF = 0;
public static final short SHN_LORESERVE = (short) 0xff00;
public static final short SHN_LOPROC = (short) 0xff00;
public static final short SHN_HIPROC = (short) 0xff1f;
public static final short SHN_LOOS = (short) 0xff20;
public static final short SHN_HIOS = (short) 0xff3f;
public static final short SHN_ABS = (short) 0xffff1;
public static final short SHN_COMMON = (short) 0xffff2;
public static final short SHN_XINDEX = (short) 0xffff;
public static final short SHN_HIRESERVE = (short) 0xffff;

/**
 * Flag informing that this section contains data that should be writable
 during process execution.
 */
public static final int FLAG_WRITE = 0x1;
/**
 * Flag informing that section occupies memory during process execution.
 */
public static final int FLAG_ALLOC = 0x2;
/**
 * Flag informing that section contains executable machine instructions.
 */
public static final int FLAG_EXEC_INSTR = 0x4;
/**
 * Flag informing that all the bits in the mask are reserved for processor
 specific semantics.
 */
public static final int FLAG_MASK = 0xf0000000;

/**
 * Name for the section containing the string table.
 * <p>
 * This section contains a string table which contains names for symbol
 structures
 * by being indexed by the {@link ElfSymbol#st_name} field.
 */
public static final String NAME_STRTAB = ".strtab";

```



```

    * Name for the section containing the dynamic string table.
    */
    public static final String NAME_DYNSTR = ".dynstr";
    /**
    * Name for the section containing read-only initialized data.
    */
    public static final String NAME_RODATA = ".rodata";

    /**
    * Index into the section header string table which gives the name of the
    section.
    */
    public final int name_ndx; // Elf32_Word or Elf64_Word - 4 bytes in both.
    /**
    * Section content and semantics.
    */
    public final int type; // Elf32_Word or Elf64_Word - 4 bytes in both.
    /**
    * Flags.
    */
    public final long flags; // Elf32_Word or Elf64_Word.
    /**
    * sh_addr. If the section will be in the memory image of a process this will
    be the address at which the first byte
    * of section will be loaded. Otherwise, this value is 0.
    */
    public final long address; // Elf32_Addr
    /**
    * Offset from beginning of file to first byte of the section.
    */
    public final long section_offset; // Elf32_Off
    /**
    * Size in bytes of the section. TYPE_NOBITS is a special case.
    */
    public final /* uint32_t */ long size;
    /**
    * Section header table index link.
    */
    public final /* uint32_t */ int link;
    /**
    * Extra information determined by the section type.
    */
    public final /* uint32_t */ int info;
    /**
    * Address alignment constraints for the section.
    */
    public final /* uint32_t */ long address_alignment;
    /**
    * Size of a fixed-size entry, 0 if none.
    */
    public final long entry_size; // Elf32_Word

    private final ElfFile elfHeader;

    /**
    * Reads the section header information located at offset.
    */
    ElfSectionHeader(final ElfParser parser, long offset) {
        this.elfHeader = parser.elfFile;
    }

```

```

        parser.seek(offset);

        name_ndx = parser.readInt();
        type = parser.readInt();
        flags = parser.readIntOrLong();
        address = parser.readIntOrLong();
        section_offset = parser.readIntOrLong();
        size = parser.readIntOrLong();
        link = parser.readInt();
        info = parser.readInt();
        address_alignment = parser.readIntOrLong();
        entry_size = parser.readIntOrLong();
    }

    /**
     * Returns the name of the section or null if the section has no name.
     */
    public String getName() {
        if (name_ndx == 0) return null;
        ElfStringTable tbl = elfHeader.getSectionNameStringTable();
        return tbl.get(name_ndx);
    }

    @Override
    public String toString() {
        return "ElfSectionHeader[name=" + getName() + ", type=0x" +
            Long.toHexString(type) + "]\n";
    }
}

```

src/net/fornwall/jelf/ElfSegment.java

```

package net.fornwall.jelf;

import java.io.IOException;

/**
 * Class corresponding to the Elf32_Phdr/Elf64_Phdr struct.
 *
 * An executable or shared object file's program header table is an array of
 * structures, each describing a segment or
 * other information the system needs to prepare the program for execution. An
 * object file segment contains one or more
 * sections. Program headers are meaningful only for executable and shared object
 * files. A file specifies its own
 * program header size with the ELF header's {@link ElfFile#ph_entry_size
 * e_phentsize} and {@link ElfFile#num_ph
 * e_phnum} members.
 *
 * http://www.sco.com/developers/gabi/latest/ch5.pheader.html#p\_type
 * http://stackoverflow.com/questions/22612735/how-can-i-find-the-dynamic-libraries-required-by-an-elf-binary-in-c
 */
public class ElfSegment {

    /** Type defining that the array element is unused. Other member values are
     * undefined. */
    public static final int PT_NULL = 0;
}

```

```

    /** Type defining that the array element specifies a loadable segment. */
    public static final int PT_LOAD = 1;
    /** The array element specifies dynamic linking information. */
    public static final int PT_DYNAMIC = 2;
    /**
     * The array element specifies the location and size of a null-terminated path
     name to invoke as an interpreter.
     * Meaningful only for executable files (though it may occur for shared
     objects); it may not occur more than once in
     * a file. If it is present, it must precede any loadable segment entry.
     */
    public static final int PT_INTERP = 3;
    /** The array element specifies the location and size of auxiliary information.
    */
    public static final int PT_NOTE = 4;
    /** This segment type is reserved but has unspecified semantics. */
    public static final int PT_SHLIB = 5;
    /**
     * The array element, if present, specifies the location and size of the
     program header table itself, both in the
     * file and in the memory image of the program. This segment type may not occur
     more than once in a file.
     */
    public static final int PT_PHDR = 6;
    /** The array element specifies the Thread-Local Storage template. */
    public static final int PT_TLS = 7;

    /** Lower bound of the range reserved for operating system-specific semantics.
    */
    public static final int PT_LOOS = 0x60000000;
    /** Upper bound of the range reserved for operating system-specific semantics.
    */
    public static final int PT_HIOS = 0x6fffffff;
    /** Lower bound of the range reserved for processor-specific semantics. */
    public static final int PT_LOPROC = 0x70000000;
    /** Upper bound of the range reserved for processor-specific semantics. */
    public static final int PT_HIPROC = 0x7fffffff;

    /** Elf{32,64}_Phdr#p_type. Kind of segment this element describes. */
    public final int type; // Elf32_Word/Elf64_Word - 4 bytes in both.
    /** Elf{32,64}_Phdr#p_offset. File offset at which the first byte of the
    segment resides. */
    public final long offset; // Elf32_Off/Elf64_Off - 4 or 8 bytes.
    /** Elf{32,64}_Phdr#p_vaddr. Virtual address at which the first byte of the
    segment resides in memory. */
    public final long virtual_address; // Elf32_Addr/Elf64_Addr - 4 or 8 bytes.
    /** Reserved for the physical address of the segment on systems where physical
    addressing is relevant. */
    public final long physical_address; // Elf32_addr/Elf64_Addr - 4 or 8 bytes.

    /** Elf{32,64}_Phdr#p_filesz. File image size of segment in bytes, may be 0. */
    public final long file_size; // Elf32_Word/Elf64_Xword -
    /** Elf{32,64}_Phdr#p_memsz. Memory image size of segment in bytes, may be 0.
    */
    public final long mem_size; // Elf32_Word
    /**
     * Flags relevant to this segment. Values for flags are defined in
     ELFSectionHeader.
     */

```

```

public final int flags; // Elf32_Word
public final long alignment; // Elf32_Word

private MemoizedObject<String> ptInterpreter;

ElfSegment(final ElfParser parser, long offset) {
    parser.seek(offset);
    if (parser.elfFile.objectSize == ElfFile.CLASS_32) {
        // typedef struct {
        //     Elf32_Word p_type;
        //     Elf32_Off p_offset;
        //     Elf32_Addr p_vaddr;
        //     Elf32_Addr p_paddr;
        //     Elf32_Word p_filesz;
        //     Elf32_Word p_memsz;
        //     Elf32_Word p_flags;
        //     Elf32_Word p_align;
        // } Elf32_Phdr;
        type = parser.readInt();
        this.offset = parser.readInt();
        virtual_address = parser.readInt();
        physical_address = parser.readInt();
        file_size = parser.readInt();
        mem_size = parser.readInt();
        flags = parser.readInt();
        alignment = parser.readInt();
    } else {
        // typedef struct {
        //     Elf64_Word p_type;
        //     Elf64_Word p_flags;
        //     Elf64_Off p_offset;
        //     Elf64_Addr p_vaddr;
        //     Elf64_Addr p_paddr;
        //     Elf64_Xword p_filesz;
        //     Elf64_Xword p_memsz;
        //     Elf64_Xword p_align;
        // } Elf64_Phdr;
        type = parser.readInt();
        flags = parser.readInt();
        this.offset = parser.readLong();
        virtual_address = parser.readLong();
        physical_address = parser.readLong();
        file_size = parser.readLong();
        mem_size = parser.readLong();
        alignment = parser.readLong();
    }

    switch (type) {
    case PT_INTERP:
        ptInterpreter = new MemoizedObject<String>() {
            @Override
            protected String computeValue() throws ElfException {
                parser.seek(ElfSegment.this.offset);
                StringBuilder buffer = new StringBuilder();
                int b;
                while ((b = parser.readUnsignedByte()) != 0)
                    buffer.append((char) b);
                return buffer.toString();
            }
        }
    }
}

```

```

        };
        break;
    }
}

@Override
public String toString() {
    String typeString;
    switch (type) {
        case PT_NULL:
            typeString = "PT_NULL";
            break;
        case PT_LOAD:
            typeString = "PT_LOAD";
            break;
        case PT_DYNAMIC:
            typeString = "PT_DYNAMIC";
            break;
        case PT_INTERP:
            typeString = "PT_INTERP";
            break;
        case PT_NOTE:
            typeString = "PT_NOTE";
            break;
        case PT_SHLIB:
            typeString = "PT_SHLIB";
            break;
        case PT_PHDR:
            typeString = "PT_PHDR";
            break;
        default:
            typeString = "0x" + Long.toHexString(type);
            break;
    }

    String pFlagsString = "";
    if (isReadable()) pFlagsString += (pFlagsString.isEmpty() ? "" : "|") +
"read";
    if (isWritable()) pFlagsString += (pFlagsString.isEmpty() ? "" : "|") +
"write";
    if (isExecutable()) pFlagsString += (pFlagsString.isEmpty() ? "" : "|") +
"execute";

    if (pFlagsString.isEmpty()) pFlagsString = "0x" + Long.toHexString(flags);

    return "ElfProgramHeader[p_type=" + typeString + ", p_filesz=" + file_size +
", p_memsz=" + mem_size + ", p_flags=" + pFlagsString + ", p_align="
+ alignment + ", range=[0x" + Long.toHexString(virtual_address) + "-
0x" + Long.toHexString(virtual_address + mem_size) + "]]";
}

/** Only for {@link #PT_INTERP} headers. */
public String getInterpreter() throws IOException {
    return (ptInterpreter == null) ? null : ptInterpreter.getValue();
}

public boolean isReadable() {
    return (flags & /* PF_R= */4) != 0;
}

```

```

    public boolean isWritable() {
        return (flags & /* PF_W= */2) != 0;
    }

    public boolean isExecutable() {
        return (flags & /* PF_X= */1) != 0;
    }
}

```

src/net/fornwall/jelf/ElfStringTable.java

```

package net.fornwall.jelf;

import java.io.IOException;

/**
 * String table sections hold null-terminated character sequences, commonly called
 * strings.
 *
 * The object file uses these strings to represent symbol and section names.
 *
 * You reference a string as an index into the string table section.
 */
final public class ElfStringTable extends ElfSection {

    /** The string table data. */
    private final byte[] data;
    public final int numStrings;

    /** Reads all the strings from [offset, length]. */
    ElfStringTable(ElfParser parser, long offset, int length, ElfSectionHeader
header) throws ElfException {
        super(parser, header);

        parser.seek(offset);
        data = new byte[length];
        int bytesRead = parser.read(data);
        if (bytesRead != length)
            throw new ElfException("Error reading string table (read " + bytesRead +
"bytes - expected to " + "read " + data.length + "bytes)");

        int stringsCount = 0;
        for (byte datum : data) if (datum == '\0') stringsCount++;
        numStrings = stringsCount;
    }

    public String get(int index) {
        int endPtr = index;
        while (data[endPtr] != '\0')
            endPtr++;
        return new String(data, index, endPtr - index);
    }
}

```

src/net/fornwall/jelf/ElfSymbol.java

```

package net.fornwall.jelf;

/**

```

** An entry in the {@Link ElfSymbolTableSection}, which holds information needed to locate and relocate a program's symbolic definitions and references.*

** <p>*

** In the elf.h header file the struct definitions are:*

** <pre>*

```
* typedef struct {  
*     uint32_t      st_name;  
*     Elf32_Addr    st_value;  
*     uint32_t      st_size;  
*     unsigned char st_info;  
*     unsigned char st_other;  
*     uint16_t      st_shndx;  
* } Elf32_Sym;
```

```
* typedef struct {  
*     uint32_t      st_name;  
*     unsigned char st_info;  
*     unsigned char st_other;  
*     uint16_t      st_shndx;  
*     Elf64_Addr    st_value;  
*     uint64_t      st_size;  
* } Elf64_Sym;
```

** </pre>*

**/*

public final class ElfSymbol {

public enum Visibility {

/**

** The visibility of symbols with the STV_DEFAULT attribute is as specified by the symbol's binding type.*

** <p>*

** That is, global and weak symbols are visible outside of their defining component, the executable file or shared object.*

** Local symbols are hidden. Global and weak symbols can also be preempted, that is, they may be interposed by definitions*

** of the same name in another component.*

**/*

STV_DEFAULT,

/**

** This visibility attribute is currently reserved.*

**/*

STV_INTERNAL,

/**

** A symbol defined in the current component is hidden if its name is not visible to other components. Such a symbol is necessarily protected.*

** <p>*

** This attribute is used to control the external interface of a component. An object named by such a symbol may still be referenced from another component if its address is passed outside.*

** <p>*

** A hidden symbol contained in a relocatable object is either removed or converted to STB_LOCAL binding by the link-editor when the relocatable object is included in an executable file or shared object.*

**/*

STV_HIDDEN,

/**

** A symbol defined in the current component is protected if it is visible in other components but cannot be preempted.*

```

    *
    * Any reference to such a symbol from within the defining component must
    be resolved to the definition in that component, even if there is a definition in
    another component that would interpose by the default rules. A symbol with
    STB_LOCAL binding will not have STV_PROTECTED visibility.
    */

```

```

    STV_PROTECTED
}

```

```

/**
 * Binding specifying that local symbols are not visible outside the object
 file that contains its definition.
 */

```

```

public static final int BINDING_LOCAL = 0;

```

```

/**
 * Binding specifying that global symbols are visible to all object files
 being combined.
 */

```

```

public static final int BINDING_GLOBAL = 1;

```

```

/**
 * Binding specifying that the symbol resembles a global symbol, but has a
 lower precedence.
 */

```

```

public static final int BINDING_WEAK = 2;

```

```

/**
 * Lower bound binding values reserved for processor specific semantics.
 */

```

```

public static final int BINDING_LOPROC = 13;

```

```

/**
 * Upper bound binding values reserved for processor specific semantics.
 */

```

```

public static final int BINDING_HIPROC = 15;

```

```

/**
 * Type specifying that the symbol is unspecified.
 */

```

```

public static final byte STT_NOTYPE = 0;

```

```

/**
 * Type specifying that the symbol is associated with an object.
 */

```

```

public static final byte STT_OBJECT = 1;

```

```

/**
 * Type specifying that the symbol is associated with a function or other
 executable code.
 */

```

```

public static final byte STT_FUNC = 2;

```

```

/**
 * Type specifying that the symbol is associated with a section. Symbol table
 entries of this type exist for

```

```

 * relocation and normally have the binding BINDING_LOCAL.
 */

```

```

public static final byte STT_SECTION = 3;

```

```

/**
 * Type defining that the symbol is associated with a file.
 */

```

```

public static final byte STT_FILE = 4;

```

```

/**
 * The symbol labels an uninitialized common block.
 */

```



```

public static final byte STT_COMMON = 5;
/**
 * The symbol specifies a Thread-Local Storage entity.
 */
public static final byte STT_TLS = 6;

/**
 * Lower bound for range reserved for operating system-specific semantics.
 */
public static final byte STT_LOOS = 10;
/**
 * Upper bound for range reserved for operating system-specific semantics.
 */
public static final byte STT_HIOS = 12;
/**
 * Lower bound for range reserved for processor-specific semantics.
 */
public static final byte STT_LOPROC = 13;
/**
 * Upper bound for range reserved for processor-specific semantics.
 */
public static final byte STT_HIPROC = 15;

/**
 * Index into the symbol string table that holds the character representation
of the symbols. 0 means the symbol has
 * no character name.
 */
public final int st_name; // Elf32_Word
/**
 * Value of the associated symbol. This may be a relative address for .so or
absolute address for other ELFs.
 */
public final long st_value; // Elf32_Addr
/**
 * Size of the symbol. 0 if the symbol has no size or the size is unknown.
 */
public final long st_size; // Elf32_Word
/**
 * Specifies the symbol type and binding attributes.
 */
public final short st_info; // unsigned char
/**
 * Currently holds the value of 0 and has no meaning.
 */
public final short st_other; // unsigned char
/**
 * Index to the associated section header. This value will need to be read as
an unsigned short if we compare it to
 * ELFSectionHeader.NDX_LORESERVE and ELFSectionHeader.NDX_HIRESERVE.
 */
public final /* Elf32_Half */ short st_shndx;

public final int section_type;

/**
 * Offset from the beginning of the file to this symbol.
 */
public final long offset;

```

```

private final ElfFile elfHeader;

ElfSymbol(ElfParser parser, long offset, int section_type) {
    this.elfHeader = parser.elfFile;
    parser.seek(offset);
    this.offset = offset;
    if (parser.elfFile.objectSize == ElfFile.CLASS_32) {
        st_name = parser.readInt();
        st_value = parser.readInt();
        st_size = parser.readInt();
        st_info = parser.readUnsignedByte();
        st_other = parser.readUnsignedByte();
        st_shndx = parser.readShort();
    } else {
        st_name = parser.readInt();
        st_info = parser.readUnsignedByte();
        st_other = parser.readUnsignedByte();
        st_shndx = parser.readShort();
        st_value = parser.readLong();
        st_size = parser.readLong();
    }

    this.section_type = section_type;

    switch (getType()) {
        case STT_NOTYPE:
            break;
        case STT_OBJECT:
            break;
        case STT_FUNC:
            break;
        case STT_SECTION:
            break;
        case STT_FILE:
            break;
        case STT_LOPROC:
            break;
        case STT_HIPROC:
            break;
        default:
            break;
    }
}

/**
 * Returns the binding for this symbol.
 */
public int getBinding() {
    return st_info >> 4;
}

/**
 * Returns the symbol type.
 */
public int getType() {
    return st_info & 0x0F;
}

```

```

/**
 * Returns the name of the symbol or null if the symbol has no name.
 */
public String getName() throws ElfException {
    // Check to make sure this symbol has a name.
    if (st_name == 0) return null;

    // Retrieve the name of the symbol from the correct string table.
    String symbol_name = null;
    if (section_type == ElfSectionHeader.SHT_SYMTAB) {
        symbol_name = elfHeader.getStringTable().get(st_name);
    } else if (section_type == ElfSectionHeader.SHT_DYNSYM) {
        symbol_name = elfHeader.getDynamicStringTable().get(st_name);
    }
    return symbol_name;
}

public Visibility getVisibility() {
    if (st_other < 0 || st_other > 3) throw new ElfException("Unsupported
st_other=" + st_other);
    return Visibility.values()[st_other];
}

@Override
public String toString() {
    String typeString;
    int typeInt = getType();
    switch (typeInt) {
        case STT_NOTYPE:
            typeString = "unspecified";
            break;
        case STT_OBJECT:
            typeString = "object";
            break;
        case STT_FUNC:
            typeString = "function";
            break;
        case STT_SECTION:
            typeString = "section";
            break;
        case STT_FILE:
            typeString = "file";
            break;
        case STT_LOPROC:
            typeString = "loproc";
            break;
        case STT_HIPROC:
            typeString = "hiproc";
            break;
        default:
            typeString = Integer.toString(typeInt);
            break;
    }

    return "ElfSymbol[name=" + getName() + ", type=" + typeString + ", size="
+ st_size + "];"
}
}

```

src/net/fornwall/jelf/ElfSymbolTableSection.java

```

package net.fornwall.jelf;

/**
 * An ELF section with symbol information.
 *
 * This class represents either of two section types:
 *
 * <ul>
 * <li>{@link ElfSectionHeader#SHT_DYNSYM}: For a minimal set of symbols
adequate for dynamic linking. Can be stripped and has no runtime cost (is non-
allocable). Normally named ".dynsym".</li>
 * <li>{@link ElfSectionHeader#SHT_SYMTAB}: A complete symbol table typically
used for link editing. Can not be stripped (is allocable). Normally named
".symtab".</li>
 * </ul>
 */
public class ElfSymbolTableSection extends ElfSection {

    public final ElfSymbol[] symbols;

    public ElfSymbolTableSection(ElfParser parser, ElfSectionHeader header) {
        super(parser, header);

        int num_entries = (int) (header.size / header.entry_size);
        symbols = new ElfSymbol[num_entries];
        for (int i = 0; i < num_entries; i++) {
            final long symbolOffset = header.section_offset + (i *
header.entry_size);
            symbols[i] = new ElfSymbol(parser, symbolOffset, header.type);
        }
    }
}

```

src/net/fornwall/jelf/MemorizedObject.java

```

package net.fornwall.jelf;

import java.io.IOException;

/**
 * A memoized object. Override {@link #computeValue} in subclasses; call {@link
#getValue} in using code.
 */
abstract class MemorizedObject<T> {
    private boolean computed;
    private T value;

    /**
     * Should compute the value of this memoized object. This will only be called
once, upon the first call to
     * {@link #getValue}.
     */
    protected abstract T computeValue() throws ElfException;

    /** Public accessor for the memoized value. */
    public final T getValue() throws ElfException {
        if (!computed) {
            value = computeValue();
            computed = true;
        }
    }
}

```

```
        return value;
    }

    @SuppressWarnings("unchecked")
    public static <T> MemoizedObject<T>[] uncheckedArray(int size) {
        return new MemoizedObject[size];
    }
}
```